Arithmetic-Geometric Structure of Agrawal's Conjecture and the Torus Dimension Count Behind the Double Condition

(AG track working document — v1.0: items AG1–AG6 complete)

Dan Shumow Claude (Anthropic Fable 5)

October 21, 2026

Abstract

We reframe Agrawal's conjecture and the Lenstra–Pomerance counterexample mechanism in the language of one-dimensional algebraic groups. The main new result (item AG2 of `doc/arith-geom-plan.md`) upgrades the empirical finding A11 of this project (`data/FINDINGS.md`) to a structural statement: for *every* AKS modulus r, the constructive (Lenstra–Pomerance-style) route forces smoothness control on the $\mathbb{F}_p$-point orders of a product of cyclotomic tori of total dimension at least 2, and the minimum 2 is attained exactly by the split and nonsplit one-dimensional tori — orders $p-1$ and $p+1$ — i.e. by the Carmichael + Lucas–Carmichael double condition (equivalently, twin-smoothness). The key identity is an inversion collapse: when $d = \mathrm{ord}_r(p)$ is even, the unique involution of $(\mathbb{Z}/r)^\times$ is -1, forcing $(X-1)^{p^{d/2}-1} \equiv -X^{-1} \pmod{p, \Phi_r(X)}$ and collapsing the order of $X-1$ into the cyclotomic levels $e \mid d/2$. Since $\sum_{e|m} \varphi(e) = m$, the burden is $d/2$ for even d and d for odd d; validity of a counterexample forces $d \geq 3$, so the burden is minimized at $d = 4$ (possible iff $4 \mid r-1$), value 2, never 1. All identities and the resulting burden table are verified numerically (`code/ag2_torus_check.py`), reproducing A11's measured table for all eleven moduli tested there. Throughout the document the collapse identities, the stabilizer and independence lemmas, and the index formulas are proved unconditionally; exactly *one* Artin-type hypothesis is used — *generic sharpness*, the attainment of the support bounds (Proposition 3.5) — named wherever it enters, empirically supported (0 exceptions across 196 sampled large primes), and a theorem in the function-field mirror. The mechanism scope (Frobenius-matching constructions, Remark 2.3) is a separate, stated restriction, not a hypothesis. Section 5 (item AG5) extends the count to Popovych's two-section conjecture: the second section $X+2$ has trivial stabilizer in the Galois action — its zero -2, unlike the zero 1 of $X-1$, is not inversion-stable — so no collapse exists at any depth, the burden rises to $\min\{d : d \mid r-1,\ d \geq 3\} \geq 3$ for every modulus (4 at Popovych's own $r = 5$: $p-1$, $p+1$ and p^2+1 simultaneously smooth), and the twin-smooth regime, the only one this project found supply-feasible, does not exist for the pair. This is the first quantified evidence in either direction on Popovych's conjecture beyond search ranges. Section 6 (item AG3) organizes the generalized-Carmichael literature into a hierarchy of *impersonation cells* indexed by cyclotomic levels and Frobenius-power parities: Carmichael, Lucas–Carmichael, Williams and Howe's higher-order Carmichael numbers are the cells $(\{1\}, 0)$, $(\{2\}, 1)$, $(\{1, 2\}, (0, 1))$ and depth-2-any-power respectively, and the striking asymmetry of the record — rigid depth-2 impersonators exist at 4.4×10^{14} (Howe) while the Williams cell is empty to 2^{64} — shows that *sign splitting*, not depth, is the constructible/blocked boundary, in exact agreement with this project's N1/A10 measurements. Section 7 (item AG4) settles the elliptic axis: a canonical elliptic Agrawal congruence exists only for CM curves, and there a single *stabilizer-index theorem* — burden $= [\langle\lambda\rangle : \langle\lambda\rangle \cap U]$, where U is the unit-stabilizer of the test section — subsumes Theorems 3.6 and 5.5 and gives the same floor of 2: the would-be burden-1 cells lie inside the exception clause (they are the elliptic analogue of rigid Carmichael numbers), the burden-2 cells demand twin-smoothness again with sign-split targets, and gate AG-G2 is not passed — the

negative result extends to the CM-elliptic axis. Appendix A (item AG6) sketches the function-field mirror, where both heuristic ingredients of the analysis are theorems by Weil's RH for curves while the problem itself carries no complexity content — locating the residual difficulty over $\mathbb{Z}$ precisely in the unconditional-equidistribution gap.

1 The Frobenius-lift dictionary

This section fixes the geometric reading used throughout (item AG1). Write $T(-1, n, r)$ for the AKS congruence

$$(X-1)^n \equiv X^n - 1 \pmod{n,\ X^r - 1}, \tag{1}$$

r an odd prime not dividing n.

Proposition 1.1 (primality as a Frobenius lift)**.** *Let $A_n = (\mathbb{Z}/n)[X]/(X^r-1)$, the mod-$n$ group algebra of μ_r, and let $\sigma_n : X \mapsto X^n$ be the n-th Adams operation (well defined since $X^r = 1$). If n is prime, then the n-power map on A_n coincides with σ_n; that is, $f(X)^n \equiv f(X^n) \pmod{n,\ X^r-1}$ for* every *f. The congruence (1) is this coincidence tested on the single element $f = X - 1$.*

Proof. For n prime the n-power map is a ring endomorphism (freshman's dream) and fixes $\mathbb{Z}/n$ (Fermat), so $f(X)^n \equiv f^{(n)}(X^n) \equiv f(X^n)$, where $f^{(n)}$ raises coefficients to the n-th power. □

Before truncation, the coincidence characterizes primality:

Proposition 1.2 (the untruncated test [1, Lemma 2.1])**.** *Let a be coprime to $n \geq 2$. Then n is prime if and only if $(X+a)^n \equiv X^n + a \pmod{n}$ in $(\mathbb{Z}/n)[X]$. (If $q^k \parallel n$ for a prime $q < n$, the coefficient of X^q is $\binom{n}{q} a^{n-q} \not\equiv 0 \bmod n$.) The AKS test, and Agrawal's conjecture, concern how much of this characterization survives the truncation from the affine line to the finite subscheme μ_r.*

In λ-ring language the statement is canonical rather than ad hoc. $\mathbb{Z}[X]/(X^r-1) = \mathbb{Z}[\mathbb{Z}/r]$ carries the Adams operations $\psi^k : X \mapsto X^k$, which pairwise commute, and each ψ^p is a *Frobenius lift*: $\psi^p(f) \equiv f^p \pmod{p}$ for every f (the calculation of Proposition 1.1). By Wilkerson's criterion [12], a torsion-free ring with such a commuting family of Frobenius lifts carries a unique λ-ring structure; in Borger's absolute geometry a Λ-structure *is* descent data to $\mathbb{F}_1$ [11]. In this language:

$$n \text{ is prime} \iff \psi^n \text{ is a Frobenius lift at } n,$$

and the AKS congruence (1) is the lift condition tested on the single section $X-1$ of μ_r. Agrawal's conjecture is therefore a *quantitative descent statement*: one well-chosen section certifies the lift condition globally, up to the $\sigma_n^2 = \mathrm{id}$ exception. The *pseudofields* of Lenstra–Pomerance [3] are the operational form of the same object — finite rings equipped with a certified Frobenius datum, which their $\tilde{O}(\log^6 n)$ algorithm constructs unconditionally; an Agrawal-type conjecture asserts the datum can be certified from a single section, at $\tilde{O}(\log^3 n)$.

Lemma 1.3 (the exceptional clause)**.** *$n^2 \equiv 1 \pmod{r}$ iff σ_n, viewed as the element $n \bmod r$ of $(\mathbb{Z}/r)^\times \cong \mathrm{Gal}(\mathbb{Q}(\zeta_r)/\mathbb{Q})$, lies in $\mathrm{Gal}(\mathbb{Q}(\zeta_r)/\mathbb{Q}(\zeta_r)^+) = \{1, \text{complex conjugation}\}$.*

Proof. $(\mathbb{Z}/r)^\times$ is cyclic of even order $r-1$; its unique element of order 2 is -1, which acts on μ_r as inversion, i.e. as complex conjugation, the generator of $\mathrm{Gal}(\mathbb{Q}(\zeta_r)/\mathbb{Q}(\zeta_r)^+)$. □

Agrawal's conjecture thus asserts: *a composite Frobenius impostor on the single section $X-1$ must act trivially on the maximal real subfield.* A *valid counterexample* is a composite n with $T(-1, n, r)$ and $t := \mathrm{ord}_r(n) > 2$.

2 Setup: cyclotomic tori and the mechanism

2.1 Cyclotomic tori

For $e \geq 1$ let Φ_e denote the e-th cyclotomic polynomial, so that

$$p^d - 1 \;=\; \prod_{e \mid d} \Phi_e(p). \tag{2}$$

Over $\mathbb{F}_p$, the Weil restriction $\mathrm{Res}_{\mathbb{F}_{p^d}/\mathbb{F}_p} \mathbb{G}_m$ decomposes up to isogeny as $\prod_{e|d} T_e$, where T_e is the cyclotomic torus of dimension $\varphi(e)$ split by $\mathbb{F}_{p^e}$, with $\#T_e(\mathbb{F}_p) = \Phi_e(p)$ [8, 9]. The two *one-dimensional* tori are $T_1 = \mathbb{G}_m$ ($\#T_1(\mathbb{F}_p) = p-1$, the split torus) and T_2 the norm-one torus of $\mathbb{F}_{p^2}/\mathbb{F}_p$ ($\#T_2(\mathbb{F}_p) = p+1$, the nonsplit twist); $\varphi(e) = 1$ iff $e \in \{1, 2\}$. We refer to the index e in (2) as the *level*.

2.2 The order of $X - 1$

Fix primes $p \neq r$, $p \nmid n$, and let $d = d_p := \mathrm{ord}_r(p)$. Then

$$\mathbb{F}_p[X]/(X^r - 1) \;\cong\; \mathbb{F}_p \times \mathbb{F}_p[X]/(\Phi_r(X)), \qquad \mathbb{F}_p[X]/(\Phi_r(X)) \;\cong\; \prod_{i=1}^{(r-1)/d} \mathbb{F}_{p^d},$$

the factors indexed by the orbits of $\langle p \rangle$ on $(\mathbb{Z}/r)^\times$; in the i-th factor X maps to a primitive r-th root of unity ζ_i. The element $X - 1$ vanishes in the $\mathbb{F}_p$ factor (where (1) holds trivially: both sides are 0) and is a unit in $\mathbb{F}_p[X]/\Phi_r$ since $\Phi_r(1) = r \neq 0$ in $\mathbb{F}_p$.

Definition 2.1. *$\rho_r(p)$ is the multiplicative order of the image of $X - 1$ in $\big(\mathbb{F}_p[X]/\Phi_r(X)\big)^\times$, i.e. $\rho_r(p) = \mathrm{lcm}_i \,\mathrm{ord}(\zeta_i - 1)$.*

For $r = 5$ this is the quantity $\rho(p)$ of Váňa [7], with $\rho(p) \mid 10(p^2 - 1)$; Corollary 3.3 below recovers this.

2.3 The Lenstra–Pomerance mechanism

Definition 2.2 (L–P mechanism)**.** *A squarefree composite $n = p_1 \cdots p_k$ coprime to r is produced by the* L–P mechanism *if for every i there is $j_i \geq 0$ with*

$$n \;\equiv\; p_i^{j_i} \pmod{\mathrm{lcm}\big(r,\ \rho_r(p_i)\big)}, \tag{3}$$

and $\mathrm{ord}_r(n) > 2$ (validity).

The original $r = 5$ construction [2, 4] is the instance $j_i = 1$: the conditions $p \equiv 3 \pmod{80}$, $p - 1 \mid n - 1$, $p + 1 \mid n + 1$ give $n \equiv p \pmod{\mathrm{lcm}(5, \rho_5(p))}$ once $\rho_5(p) \mid 10(p^2 - 1)$ is known. Our entire Track A pipeline (`PLAN.md`) is a search inside Definition 2.2.

Remark 2.3 (scope). The mechanism is sufficient, not necessary, for (1): the component equation $(\zeta - 1)^n = \zeta^n - 1$ could conceivably hold "accidentally" for some p, outside the Frobenius-matching family (3). No such solutions are known, and no way to construct them is known; every published counterexample construction [2, 4, 6] is an instance of Definition 2.2. All results below quantify *this mechanism*, in line with the scope discipline of the negative-result paper (`doc/negative-result-plan.md`).

3 The torus dimension count

3.1 Frobenius matching and the inversion collapse

Lemma 3.1 (Frobenius matching). *With notation as above:*

(i) $(X-1)^{p^k} \equiv X^{p^k} - 1 \pmod{p,\, X^r - 1}$ *for all* $k \geq 0$.

(ii) If $n \equiv p^j \pmod{\operatorname{lcm}(r, \rho_r(p))}$ *then* $T(-1, n, r)$ *holds modulo* p.

(iii) Conversely, (3) *forces* $n \equiv p^j \pmod{r}$, *hence* $\operatorname{ord}_r(n) \mid d_p$. *In particular a valid counterexample (*$\operatorname{ord}_r(n) > 2$*) forces* $d_p \geq 3$ *for every* $p \mid n$.

Proof. (i) is the iterated Frobenius endomorphism of $\mathbb{F}_p$-algebras. (ii): write $n = p^j + c \cdot \operatorname{lcm}(r, \rho_r(p))$. In the $\mathbb{F}_p$ factor both sides of (1) vanish. In $\mathbb{F}_p[X]/\Phi_r$: $(X-1)^n = (X-1)^{p^j} \left((X-1)^{\rho_r(p)}\right)^{c'} = (X-1)^{p^j} \equiv X^{p^j} - 1$ by (i), and $X^{p^j} = X^{n \bmod r} = X^n$ since $X^r \equiv 1 \bmod \Phi_r$ (as $X^r - 1 = (X-1)\Phi_r(X)$ for r prime). (iii): reduce (3) modulo r; then $n \in \langle p \rangle \subseteq (\mathbb{Z}/r)^\times$, so $\operatorname{ord}_r(n)$ divides $\#\langle p \rangle = d_p$, and $\operatorname{ord}_r(n) \geq 3$ gives $d_p \geq 3$. □

Lemma 3.2 (inversion collapse). *Suppose* $d = \operatorname{ord}_r(p)$ *is even. Then, in* $\mathbb{F}_p[X]/\Phi_r(X)$,

$$(X-1)^{p^{d/2}-1} \;=\; -X^{-1},$$

and consequently

$$\rho_r(p) \;\Big|\; 2r\,(p^{d/2} - 1) \;=\; 2r \prod_{e \mid d/2} \Phi_e(p). \tag{4}$$

Proof. $(\mathbb{Z}/r)^\times$ is cyclic, so its unique element of order 2 is -1; since $\operatorname{ord}_r(p) = d$, the element $p^{d/2}$ has order 2, whence $p^{d/2} \equiv -1 \pmod{r}$. In each field component, with $X \mapsto \zeta$ a primitive r-th root of unity, iterated Frobenius gives

$$(\zeta - 1)^{p^{d/2}} = \zeta^{p^{d/2}} - 1 = \zeta^{-1} - 1 = -\zeta^{-1}(\zeta - 1),$$

and dividing by $\zeta - 1 \neq 0$ yields $(\zeta - 1)^{p^{d/2}-1} = -\zeta^{-1}$. This holds in every component, hence in the product ring. Finally $(-X^{-1})^{2r} = X^{-2r} = 1$, so $(X-1)^{2r(p^{d/2}-1)} = 1$, giving (4); the product form is (2) applied to $p^{d/2} - 1$. □

Corollary 3.3 (Váňa's bound). *For* $r = 5$ *and* $p \equiv \pm 2 \pmod{5}$ *(so* $d = 4$*),* $\rho_5(p) \mid 10(p^2 - 1)$.

Definition 3.4 (support level). *Set*

$$d' \;=\; d'(p) \;:=\; \begin{cases} d/2, & d \text{ even}, \\ d, & d \text{ odd}. \end{cases}$$

For odd d the trivial bound $\rho_r(p) \mid p^d - 1 = \prod_{e \mid d} \Phi_e(p)$ applies (orders in a product of copies of $\mathbb{F}_{p^d}^\times$), so in both cases

$$\rho_r(p) \;\Big|\; 2r \prod_{e \mid d'} \Phi_e(p)\,: \tag{5}$$

the p-dependent part of $\rho_r(p)$ *is supported on the cyclotomic levels* $e \mid d'$.

3.2 Generic sharpness

The bound (5) is an upper bound on support; the mechanism's cost analysis also needs it to be generically *attained* — otherwise a harvest could fish for primes p where $\rho_r(p)$ collapses further.

Proposition 3.5 (generic sharpness; heuristic + empirical)**.** *Let $\ell > 2rd$ be a prime with $\ell \mid \Phi_e(p)$ for some $e \mid d'$. Then $\ell \mid \rho_r(p)$ unless in every component the image of $X-1$ is an ℓ-th power in $\mathbb{F}_{p^d}^\times$. Modelling the components as independent uniform elements, the exceptional probability is $\ell^{-(r-1)/d} \le 1/\ell$. Consequently, for all p outside a set of (heuristic) density $O(\sum_\ell 1/\ell)$, the large-prime part of $\prod_{e\mid d'} \Phi_e(p)$ divides $\rho_r(p)$, and requiring "$\rho_r(p)$ is B-smooth" is equivalent to requiring "$\Phi_e(p)$ is B-smooth for every $e \mid d'$".*

Evidence. $\ell \nmid \rho_r(p)$ iff the image of $X-1$ has trivial ℓ-part in each cyclic component group $\mathbb{F}_{p^d}^\times$, i.e. is an ℓ-th power in every component whose order ℓ divides; the ℓ-th powers have index ℓ in each component containing ℓ-torsion. Rigorously controlling the image of $\zeta - 1$ is an open Artin-type problem; empirically (`code/ag2_torus_check.py`, 67 sample primes across $r \in \{5,7,11,13,19,23,41\}$, $d \in \{3,4,5,6,11\}$): all 37 tested large primes ℓ divided $\rho_r(p)$ (0 exceptions; ≈ 0.19 expected under the $1/\ell$ model), and in all 67 cases removing the top level $e = d'$ from the bound (5) broke it (the top level is genuinely present in $\rho_r(p)$). □

3.3 The dimension-count theorem

Theorem 3.6 (torus dimension count)**.** *Let r be an odd prime and let n be produced by the L–P mechanism (Definition 2.2) with $\mathrm{ord}_r(n) > 2$. For each $p \mid n$ put $d_p = \mathrm{ord}_r(p)$ and d_p' as in Definition 3.4. Assume generic sharpness (Proposition 3.5) at p. Then, for the CRT system (3) to admit assembly over many primes with a common B-smooth modulus (the setting of the Track-A pipeline), each $p \mid n$ must satisfy:*

$$\Phi_e(p) \text{ is } B\text{-smooth for every } e \mid d_p' \quad\Longleftrightarrow\quad \#\Big(\textstyle\prod_{e\mid d_p'} T_e\Big)(\mathbb{F}_p) \text{ is } B\text{-smooth},$$

a smoothness condition on the $\mathbb{F}_p$-points of a torus of dimension

$$\dim \prod_{e\mid d_p'} T_e \;=\; \sum_{e\mid d_p'} \varphi(e) \;=\; d_p' \;=\; \begin{cases} d_p/2, & d_p \text{ even}, \\ d_p, & d_p \text{ odd}. \end{cases}$$

Moreover:

(a) $d_p' \ge 2$ *for every $p \mid n$: the burden is never a single one-dimensional torus. (Indeed $d_p' = 1$ would force $d_p \in \{1,2\}$, contradicting $d_p \ge 3$ from Lemma 3.1(iii).)*

(b) $d_p' = 2$ *iff $d_p = 4$, in which case the level set is $\{1,2\}$ and the condition is: $p-1 = \#T_1(\mathbb{F}_p)$ and $p+1 = \#T_2(\mathbb{F}_p)$ both B-smooth — the Carmichael + Lucas–Carmichael* double condition, *i.e. twin-smoothness of $\big(\frac{p-1}{2}, \frac{p+1}{2}\big)$.*

(c) $d_p = 4$ *is realizable iff $4 \mid r-1$; hence for $r \equiv 1 \pmod 4$ the minimal burden is 2, identical in structure to the $r = 5$ case, while for $r \equiv 3 \pmod 4$ every admissible d_p gives burden ≥ 3.*

(d) The minimal burden as a function of r is

$$\mathrm{burden}(r) \;=\; \min_{\substack{d\mid r-1 \\ d\ge 3}} \begin{cases} d/2, & d \text{ even} \\ d, & d \text{ odd} \end{cases}$$

which reproduces the measured A11 table for all eleven moduli tested there (Section 4).

Proof. The support statement is (5) plus Proposition 3.5; the dimension computation is $\sum_{e|m} \varphi(e) = m$ with $m = d'_p$, together with $\dim T_e = \varphi(e)$ and $\#T_e(\mathbb{F}_p) = \Phi_e(p)$ (Section 2). (a): $d_p \geq 3$ by Lemma 3.1(iii); $d'_p = 1$ requires $d_p \leq 2$. (b): $d'_p = 2$ with d_p even means $d_p = 4$ and levels $\{e : e \mid 2\} = \{1, 2\}$; $d'_p = 2$ with d_p odd is impossible ($d_p = 2$ is even). With p odd, B-smoothness of $p \mp 1 = 2m, 2(m+1)$ is twin-smoothness of the pair $(m, m+1)$ up to the factor 2. (c): $d_p = 4$ requires $4 \mid \#(\mathbb{Z}/r)^\times = r - 1$; conversely if $4 \mid r - 1$ the cyclic group $(\mathbb{Z}/r)^\times$ has elements of order 4, and primes p in those residue classes exist (Dirichlet), with $\mathrm{ord}_r(n) = 4 \mid d_p$ satisfiable. For $r \equiv 3 \pmod 4$: even $d \mid r - 1$ satisfy $d \equiv 2 \pmod 4$, so even $d \geq 3$ means $d \geq 6$, burden $d/2 \geq 3$; odd $d \geq 3$ gives burden $d \geq 3$. (d) is the minimization of d'_p over admissible d_p. □

Corollary 3.7 (structural closure of option (c))**.** *For every AKS modulus r, the L–P mechanism requires each pool prime to satisfy a smoothness condition on the point-orders of a torus of dimension at least 2; no modulus reduces the requirement to a single one-dimensional torus (a single condition of the shape "$p \pm 1$ smooth"). The minimum, dimension exactly 2, is achieved precisely by the split/nonsplit pair $\{T_1, T_2\}$ — the twin-smooth double condition — and only for $r \equiv 1 \pmod 4$. The double condition measured empirically in A11 is therefore intrinsic: it is the unique minimizer of a dimension count, not an artifact of $r = 5$.*

Remark 3.8 (cost interpretation). Under the standard Dickman heuristic, B-smoothness of $\Phi_e(p)$, a quantity of size $p^{\varphi(e)}$, has cost governed by $u = \varphi(e) \log p / \log B$; total burden across the level set is governed by $\sum_{e|d'} \varphi(e) = d'$. Burden 2 is the twin-smooth regime this project has harvested at scale; burden ≥ 3 requires smooth values of $\Phi_e(p) \sim p^{\varphi(e)}$ with $\varphi(e) \geq 2$, the same quantities whose *non*-smoothness underpins torus-based cryptography [9]. This is the structural content of A11's "every other r is strictly worse".

Remark 3.9 (what remains conditional). Lemmas 3.1 and 3.2, the support bound (5), and parts (a)–(d) of Theorem 3.6 as statements about the support of $\rho_r(p)$ are unconditional. The conditional ingredient is generic sharpness (Proposition 3.5): a harvest could in principle target the thin set of primes where $\rho_r(p)$ collapses below (5). Such primes have heuristic density $O(1/\ell)$ per large prime ℓ avoided — i.e. the "collapsed" pool is far thinner than the twin-smooth pool itself — so this escape is uneconomical, but it is recorded here as the localized gap per gate AG-G1 (`doc/arith-geom-plan.md`).

4 Numerical validation

`code/ag2_torus_check.py` (deterministic; output in `data/generated/ag2_torus_check.json`) verifies, over 67 sample primes spanning $(r, d) \in \{(5,4), (13,4), (41,4), (7,3), (7,6), (19,3), (11,5), (23,11)\}$ — sampled independently of the A11 scan, which drew $p > 1000$ in one fixed residue class per r, whereas here p runs from 3 upward over all classes of the given order, and the implementation is independent (exact orders by polynomial arithmetic in $\mathbb{F}_p[X]/\Phi_r$ rather than A11's level attribution):

check	statement	result
V1	Frobenius matching (Lemma 3.1(i))	67/67
V2	collapse identity $(X-1)^{p^{d/2}-1} = -X^{-1}$ (Lemma 3.2)	36/36
V3	exact $\rho_r(p)$ computed within bound (5)	67/67
V4	large primes of levels $e \mid d'$ divide $\rho_r(p)$ (Prop. 3.5)	37/37
V5	top level $e = d'$ genuinely needed in $\rho_r(p)$	67/67
V6	burden formula (Thm. 3.6(d)) vs. A11 table	11/11

V4's 0 exceptions are consistent with the $1/\ell$ model (≈ 0.19 misses expected over this sample). V6 reproduces, for $r \in \{5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41\}$, both the burden value and the minimizing d of the A11 measurement (`data/FINDINGS.md`, 2026-08-24).

5 A quantified heuristic for Popovych's conjecture

Popovych's conjecture [4] strengthens the hypothesis of Agrawal's to the simultaneous pair

$$(X-1)^n \equiv X^n - 1 \quad \text{and} \quad (X+2)^n \equiv X^n + 2 \pmod{n,\, X^r - 1}, \tag{6}$$

with the same conclusion. Beyond the Primaboinca search range ($n < 10^{17}$) there is no published evidence in either direction. This section quantifies what the second section costs the L–P mechanism — the only known counterexample engine — and finds that it destroys exactly the structure that made Agrawal's conjecture attackable: the inversion collapse of Lemma 3.2, and with it the twin-smooth (burden-2) regime.

Throughout, $r \neq 3$ is prime and $p \nmid 3(2^r+1)$; the finitely many excluded p are degenerate ($p = 3$ makes $-2 \equiv 1$, so the two sections of (6) *coincide* mod 3; $p \mid (2^r+1)/3$ makes $X+2$ a zero divisor in some component).

Lemma 5.1 (the level-one component is base-3 Fermat)**.** *The $X = 1$ component of the second congruence in (6) is $3^n \equiv 3 \pmod{n}$. In particular it is automatic for any Carmichael-type n ($p - 1 \mid n - 1$ for all $p \mid n$), hence free for the L–P construction; but unlike the first section, whose $X = 1$ component is $0 = 0$, the second section is already nontrivial at level one.*

Proof. Evaluate (6) at $X = 1$: $(1+2)^n \equiv 1 + 2 \pmod{n}$. □

Lemma 5.2 (global independence)**.** *In $\mathbb{Z}[\zeta_r]$ the elements $\zeta - 1$ and $\zeta + 2$ have principal ideals with disjoint prime support: $N(\zeta - 1) = \Phi_r(1) = r$ while $N(\zeta + 2) = \Phi_r(-2) = (2^r+1)/3$, and $\gcd\big(r,\, 2^r + 1\big) = \gcd(r, 3) = 1$. Since neither is a unit, $\zeta - 1$ and $\zeta + 2$ generate a rank-2 subgroup of $\mathbb{Q}(\zeta_r)^\times$ modulo roots of unity: no multiplicative identity ties the order of one reduction to the order of the other.*

Proof. $X^r - 1 = (X-1)\Phi_r(X)$ gives $\Phi_r(-2) = ((-2)^r - 1)/(-3) = (2^r+1)/3$. By Fermat $2^r \equiv 2 \pmod{r}$, so $2^r + 1 \equiv 3 \pmod{r}$ and $\gcd(r, 2^r+1) = \gcd(r, 3) = 1$ as $r \neq 3$. Disjoint support plus non-unit norms gives independence modulo torsion. □

Lemma 5.3 (trivial stabilizer: no collapse at any depth)**.** *For every $\gamma \in (\mathbb{Z}/r)^\times$ with $\gamma \neq 1$, the ratio $\sigma_\gamma(\zeta+2)/(\zeta+2) = (\zeta^\gamma + 2)/(\zeta + 2)$ is not a root of unity — it is not even a unit of $\mathbb{Z}[\zeta_r]$. Consequently the mechanism of Lemma 3.2 has no analogue for $X+2$: no relation $(X+2)^{p^e-1} \in \mu_{2r}$ holds for any e with $p^e \not\equiv 1 \pmod{r}$.*

Proof. Let $\mathfrak{q}$ be a prime of $\mathbb{Z}[\zeta_r]$ dividing $(\zeta + 2)$, so $\zeta \equiv -2 \pmod{\mathfrak{q}}$ and the image of -2 in the residue field has multiplicative order r (it is the image of ζ). If $\mathfrak{q}$ also divided $(\zeta^\gamma + 2)$ then $\zeta^\gamma \equiv -2 \equiv \zeta$, so $\zeta^{\gamma-1} \equiv 1 \pmod{\mathfrak{q}}$, forcing $r \mid \gamma - 1$, i.e. $\gamma = 1$. So for $\gamma \neq 1$ the ratio has nonzero valuation at $\mathfrak{q}$ and is not a unit.

Contrast with the first section: $\sigma_{-1}(\zeta - 1)/(\zeta - 1) = -\zeta^{-1} \in \mu_{2r}$, available because the ideal $(\zeta - 1)$ (the totally ramified prime above r) is fixed by the whole Galois group — geometrically, the zero locus $\{1\}$ of $X - 1$ is Galois-stable, while the zero -2 of $X + 2$ has trivial stabilizer. □

Corollary 5.4 (full-level support of the second section)**.** *Define $\rho'_r(p)$ as the multiplicative order of $X+2$ in $(\mathbb{F}_p[X]/\Phi_r(X))^\times$. Then $\rho'_r(p) \mid p^d - 1 = \prod_{e\mid d} \Phi_e(p)$, and — in the generic-sharpness sense of Proposition 3.5, now justified globally by Lemma 5.3 — its support includes* all *levels $e \mid d$, in particular the top level $\Phi_d(p)$ of dimension $\varphi(d)$, for even d the level the first section never sees.*

Theorem 5.5 (Popovych burden)**.** *Run the L–P mechanism for the pair (6): for each $p \mid n$,*

$$n \equiv p^{j_p} \pmod{\operatorname{lcm}\big(r,\ \rho_r(p),\ \rho'_r(p)\big)}, \qquad \operatorname{ord}_r(n) > 2.$$

Under generic sharpness for both sections, the combined support is all levels $e \mid d_p$, so the smoothness burden at p is $\sum_{e\mid d_p} \varphi(e) = d_p$ for every *parity of d_p. Hence*

$$\operatorname{burden}_{\mathrm{Pop}}(r) \;=\; \min\{\, d : d \mid r-1,\ d \geq 3 \,\} \;\geq\; 3 \qquad \textit{for every prime } r,$$

and:

(a) the global optimum is 3 (attained iff $3 \mid r-1$), versus 2 for the single section (Theorem 3.6);

(b) at $r = 5$ — Popovych's own proposed modulus — and generally whenever $4 \mid r-1$ and $3 \nmid r-1$, the burden is 4: the pool primes must have $p-1$, $p+1$ and p^2+1 simultaneously smooth;

(c) $\operatorname{burden}_{\mathrm{Pop}}(r) \geq \operatorname{burden}_{\mathrm{Agr}}(r)$ for every r, strictly whenever $\operatorname{burden}_{\mathrm{Agr}}(r) = 2$ (i.e. for every $r \equiv 1 \pmod 4$);

(d) the twin-smooth regime (burden 2) — the only regime this project found supply-feasible (gates G1/G2) — does not exist for the pair at any modulus.

Proof. For odd d both sections are supported on all $e \mid d$ already (Section 3), total d; the pair ties with the single section there. For even d the first section collapses to $e \mid d/2$ (Lemma 3.2) but the second does not (Corollary 5.4), so the lcm spans all $e \mid d$, total $\sum_{e\mid d} \varphi(e) = d$ — the collapse is destroyed. Admissibility $d \geq 3$ is Lemma 3.1(iii), unchanged. Minimizing d over $d \mid r-1$, $d \geq 3$ gives the formula. (a): the least divisor ≥ 3 of $r-1$ is 3 iff $3 \mid r-1$. (b): if $3 \nmid r-1$ and $4 \mid r-1$ the least is 4, with levels $\{1,2,4\}$ and $\Phi_4(p) = p^2+1$. (c): $\operatorname{burden}_{\mathrm{Agr}}(r)$ minimizes d' where $d' \leq d$ always, so it is $\leq \operatorname{burden}_{\mathrm{Pop}}(r)$; when it equals 2 the Popovych value is ≥ 3 by the displayed bound. (d) is (a) plus Theorem 3.6(b). □

Proposition 5.6 (cost at the measured scale)**.** *At the frozen harvest spec of this project (`doc/harvest-spec.md`: pool primes $p \approx 51$ bits, measured supply/demand ratio 13.8, pool 2.5×10^5), the extra quadratic level of Theorem 5.5(b) is a pure Dickman tax on a ≈ 102-bit quantity: $\rho(u)$ with $u = 102/\log_2 B$, i.e.*

$$u \approx 9.9,\ \rho(u) \approx 3.9 \times 10^{-11} \quad (B = 1259), \qquad u \approx 7.5,\ \rho(u) \approx 1.8 \times 10^{-7} \quad (B' = 12590),$$

multiplying the measured ratio down to 5×10^{-10}–2.5×10^{-6} and the pool to 10^{-5}–0.05 elements: the pair-mechanism harvest is empty at any feasible scale. *Unlike $p \mp 1$, where partition-first harvesting lets one side be enumerated rather than tested (`data/FINDINGS.md`, A1), $\Phi_4(p)$ is pinned once p is, so no enumeration trick removes this tax. On the solver side the group acquires the level-4 components, so $\log_2 |G|$ grows and the identity density falls below the measured $\rho \approx 0.004$ of H7 pools: the N1/A10 obstruction for the pair strictly contains the one for the single section.*

Remark 5.7 (interpretation and scope). Theorem 5.5 and Proposition 5.6 are, to our knowledge, the first quantified evidence in either direction on Popovych's conjecture beyond exhaustive search: every known counterexample mechanism pays at least one more torus dimension against the pair than against the single section globally (3 vs. 2), and two more at the canonical $r = 5$ (4 vs. 2); the burden-2 regime that made the Agrawal attack supply-feasible does not exist for the pair. The caveats are those of Remarks 2.3 and 3.9, applied to both sections: this quantifies the failure of the only known engine, not the truth of the conjecture. Note also the char-3 degeneracy: modulo $p = 3$ the two sections coincide, so the pair constrains only $p > 3$ divisors.

Remark 5.8 (the constant 2 is immaterial). For any integer $c \notin \{0, \pm 1\}$ the argument of Lemma 5.3 runs verbatim for the section $X + c$: at a prime $\mathfrak{q} \mid (\zeta + c)$ the image of $-c$ has multiplicative order r, so no $\gamma \neq 1$ preserves the divisor, the stabilizer is trivial, and the full burden jump of Theorem 5.5 holds with c entering only through the norm $\Phi_r(-c) = (c^r + 1)/(c + 1)$ (whose primes, with $p \mid c(c + 1)$, are the finitely many excluded p), the excluded modulus $r \mid c + 1$ (for $c = 2$: $r = 3$), and the level-one component, the base-$(c+1)$ Fermat condition. What must be avoided is exactly the rational torsion $\mu_2(\mathbb{Q}) = \{\pm 1\}$ as the zero $-c$: $c = -1$ is the first section itself, and $c = +1$ — the one innocent-looking bad choice — has its zero at -1, fixed by inversion, so the collapse *survives*: $(\zeta + 1)^{p^{d/2}} = \zeta^{-1}(\zeta + 1)$, giving $\mathrm{ord}(X + 1) \mid r(p^{d/2} - 1)$ for even d; worse, $\zeta + 1 = (\zeta^2 - 1)/(\zeta - 1)$ is multiplicatively dependent on the Galois orbit of $\zeta - 1$, so the pair $\{X - 1, X + 1\}$ adds no rank and no burden at all — a Popovych-style conjecture with $c = 1$ would be no harder for the mechanism than Agrawal's. In the language of Theorem 7.3: $U(X + c) = \{\pm 1\}$ for $c = \pm 1$ and $\{1\}$ otherwise; $c = 2$ is simply the smallest valid choice ($c = -2$ works equally, with Mersenne norm $2^r - 1$). Verified numerically (`code/ag5b_constant_check.py`): at even d the section $X + 1$ collapses in 22/22 samples with the top level never needed, while $c \in \{2, 3, -2, 5\}$ collapse only at the single degenerate point $(r, p) = (13, 5)$, where $\Phi_4(5) = 26$ has no prime outside $2r$ and the top level is vacuous for every section.

Proposition 5.9 (the $X+1$ variant is mechanism-equivalent to Agrawal's conjecture)**.** *Suppose* $n \equiv p^j \pmod{\mathrm{lcm}(r, \rho_r(p))}$ *and* $p - 1 \mid n - 1$ *(the mechanism congruence together with the Carmichael condition, both imposed by every L–P-style construction). Then* $T(1, n, r)$ *holds modulo* p *outright. Consequently every counterexample to Agrawal's conjecture produced by the mechanism also defeats the variant conjecture whose hypothesis is the pair* $\{T(-1, n, r), T(1, n, r)\}$*: that variant, though formally stronger-looking, is exactly as resistant as Agrawal's — the second hypothesis is free armor — whereas* $c = 2$ *(or any* $c \notin \{0, \pm 1\}$*) genuinely blocks the engine.*

Proof. At the $X = 1$ component, $(1 + 1)^n = 2^n \equiv 2 \pmod{p}$ since $\mathrm{ord}_p(2) \mid p - 1 \mid n - 1$. In the Φ_r components, $\zeta + 1 = (\zeta^2 - 1)/(\zeta - 1)$, a unit (its norm is $\Phi_r(-1) = 1$), and both $\zeta - 1$ and $\zeta^2 - 1$ are values of $X - 1$ at primitive r-th roots of unity, so their orders divide $\rho_r(p) = \mathrm{lcm}_a \,\mathrm{ord}(\zeta^a - 1)$. Hence $(\zeta^a - 1)^n = (\zeta^a - 1)^{p^j} = \zeta^{ap^j} - 1$ for $a = 1, 2$, and

$$(\zeta + 1)^n \;=\; \frac{(\zeta^2 - 1)^n}{(\zeta - 1)^n} \;=\; \frac{\zeta^{2p^j} - 1}{\zeta^{p^j} - 1} \;=\; \zeta^{p^j} + 1 \;=\; \zeta^n + 1,$$

using $n \equiv p^j \pmod{r}$ in the last step. Measured (`code/ag5b_constant_check.py`, check W): 48/48 synthetic mechanism instances pass $T(1, n, r)$, while 0/48 pass $T(2, n, r)$. □

5.1 Numerical validation

`code/ag5_popovych_check.py` (deterministic; output in `data/generated/ag5_popovych_check.json`) verifies, over 56 sample primes spanning the same (r, d) grid as Section 4:

check	statement	result
P1	Frobenius matching for $X+2$	56/56
P2	the $X-1$ collapse bound does *not* absorb $X+2$ (even d)	35/36
P3	top level $e=d$ present in $\rho'_r(p)$	55/56
P4	large primes of *all* levels $e \mid d$ divide $\rho'_r(p)$	47/47
P5	burden table: $\text{burden}_{\text{Pop}}$ vs. $\text{burden}_{\text{Agr}}$, 11 moduli	as stated
P6	$\gcd(r, 2^r+1) = 1$ (Lemma 5.2)	11/11
P7	Dickman cost table (Proposition 5.6)	as stated

The single P2/P3 exception is $r = 13$, $p = 5$: there $\Phi_4(5) = 26 = 2 \cdot 13$ contains no prime outside $2r$ at all, so the top level is vacuous for this one tiny prime — a small-prime degeneracy, not a collapse event; it is consistent with (indeed predicted by) the generic model, under which exceptions have probability $\sim 1/\ell$ per large prime ℓ and ≈ 0.28 misses were expected over this sample. The burden deltas of P5: $+2$ at $r \in \{5, 17, 29, 41\}$, $+1$ at $r \in \{13, 37\}$, 0 at the $r \equiv 3 \pmod 4$ moduli (where the Agrawal burden was already ≥ 3 and both conjectures share full depth).

6 The impersonation hierarchy

Theorems 3.6 and 5.5 are statements about which congruence cells a counterexample mechanism forces its pool primes into. This section defines those cells once and for all, maps the classical generalized-Carmichael literature onto them, and records what is known about each — the related-work backbone and vocabulary for the negative-result paper (`doc/negative-result-plan.md`).

Definition 6.1 (impersonation cells)**.** *Let p be prime and n coprime to p. For a finite set S of levels and powers $\mathbf{i} = (i_e)_{e \in S}$, say n* impersonates Frobenius at p on S with powers $\mathbf{i}$ *if*

$$\Phi_e(p) \mid n - p^{i_e} \qquad \text{for every } e \in S.$$

By the dictionary of Section 2 this says exactly that $x^n = x^{p^{i_e}}$ for all x in the order-$\Phi_e(p)$ subgroup of $\mathbb{F}_{p^e}^\times$ — the points of the torus T_e. Uniform-power impersonation at depth m *is the single congruence $p^m - 1 \mid n - p^i$, i.e. $x^n = x^{p^i}$ for all $x \in \mathbb{F}_{p^m}^\times$. A squarefree composite n* lies in a cell *when every $p \mid n$ impersonates accordingly. Note two normalizations: at level 1 all powers coincide ($p \equiv 1 \bmod p-1$), so only the cell membership matters there; and on T_2 one has $p^i \equiv (-1)^i \pmod{p+1}$, so at level 2 only the* parity *of i_2 matters — the parity is a sign: even $\leftrightarrow$ target $+1$, odd $\leftrightarrow$ target -1.*

Proposition 6.2 (the classical objects in cell coordinates)**.** *For squarefree composite n:*

(i) n is a Carmichael number (Korselt: $p-1 \mid n-1$ for all $p \mid n$) iff n lies in the cell $(\{1\}, 0)$.

(ii) n is a Lucas–Carmichael number ($p+1 \mid n+1$) iff n lies in $(\{2\}, 1)$: indeed $p+1 \mid n-p \iff p+1 \mid n+1$.

(iii) n is a Williams number [14, 15, 16] (both conditions) iff n lies in $(\{1,2\}, (0,1))$, equivalently iff n impersonates at depth 2 with the uniform power 1: $n \equiv p \pmod{(p^2-1)/2}$ for every $p \mid n$.

(iv) n is a Carmichael number of order m in the sense of Howe [17] ($x \mapsto x^n$ an endomorphism of every $\mathbb{Z}/n$-algebra generated by m elements) iff for every $p \mid n$ and every depth $r \leq m$ some power works: $\exists i,\ p^r - 1 \mid n - p^i$ (Howe's Theorem 1). For $m = 2$ this is: Carmichael, plus level-2 membership with some *parity. Howe's* rigid *numbers are the all-powers-0 (all-even-parity) case $\operatorname{lcm}_{r \leq m}(p^r - 1) \mid n - 1$.*

(v) *The Lenstra–Pomerance pool for Agrawal's conjecture is the Williams cell refined by the local conditions ($p \equiv 3 \bmod 80$, validity $n \not\equiv \pm 1 \bmod r$); Theorem 3.6 says the AKS mechanism at any modulus r forces the level set $\{e : e \mid d'\}$, with the Williams cell $\{1, 2\}$ as the unique dimension-2 minimum; Theorem 5.5 says the Popovych pair forces the full level set $\{e : e \mid d\}$. Grantham's Frobenius pseudoprimes [22] are the same impersonation idea tested against a single polynomial section rather than all torus points.*

Proof. (i)–(iii) unwind Definition 6.1 ((iii): $n \equiv 1 \bmod p-1$ and $n \equiv -1 \equiv p \bmod p+1$ combine by CRT to $n \equiv p \bmod \operatorname{lcm}(p-1, p+1) = (p^2-1)/2$); the point-identity form of the dictionary is verified numerically in Section 6.1. (iv) is Howe's Theorem 1 [17] together with the level-1 normalization above. (v) collects Sections 3 and 5. □

Theorem 6.3 (status of the cells — survey)**.** *(a)* $(\{1\}, 0)$*, Carmichael: infinitely many [13];* 105 *below* 10^7*.*

(b) $(\{2\}, 1)$*, Lucas–Carmichael: infinitely many — Wright [19], as a by-product of the infinitude of elliptic Carmichael numbers;* 135 *below* 10^7*.*

(c) $(\{1, 2\}, (0, 1))$*, Williams: open whether the cell is even nonempty. Williams [14] posed the question and showed such n need at least four prime factors; McIntosh–Mitra [16] verified none exist below $2^{64} \approx 1.8 \times 10^{19}$ and proved finiteness results (only finitely many with a given set of all but three prime factors). The Agrawal cell (v) is a refinement, a fortiori empty there — consistent with OEIS A329223's emptiness against Pinch's tables.*

(d) Depth-2, any parity (Howe order 2): nonempty*. Howe's explicit example* $n = 443372888629441 = 17 \cdot 31 \cdot 41 \cdot 43 \cdot 89 \cdot 97 \cdot 167 \cdot 331$ *is rigid of order 2 ($p^2 - 1 \mid n - 1$ for all eight primes; re-verified here, Section 6.1). Every order-m Carmichael number has at least $m + 2$ prime factors [18]; no order-3 example is known; Howe's heuristic predicts infinitely many of every order.*

(e) Elliptic rungs: Silverman's elliptic Korselt criteria [20], conditional infinitude of elliptic Carmichael numbers by Ekstrom–Pomerance–Thakur [21], unconditional by Wright [19] — the structures item AG4 will build on.

Remark 6.4 (sign splitting, not depth, is the wall)*.* Contrast (c) and (d): rigid depth-2 impersonators exist at 4.4×10^{14}, while the Williams cell — the *same* level set — is empty to 2^{64}. Depth is not the obstruction; the sign structure is. The rigid cell's condition is one-sided: $n \equiv 1 \pmod{\operatorname{lcm}_{p \mid n}(p^2 - 1)}$ is an AGP-type problem — harvest primes with $p^2 - 1 \mid \Lambda$ for a single highly composite Λ (the $p-1$ and $p+1$ parts may share prime support freely) and solve one subset-product with target 1; that machine produced (d) at modest size. The Williams parities $(0, 1)$ instead force target $+1$ on the level-1 side and -1 on the level-2 side; on any shared prime power the two targets conflict, so the modulus must split into coprime sides — exactly the side-coprimality (working paper Lemma 2.3) whose harvest collapse this project measured as N1 and whose two-sided solver obstruction is A10. In cell language: the project's negative results concern precisely the *sign-split* cells, and Howe's example shows the unsplit cells are reachable — the empirical record (c) vs. (d) confirms the split, not the depth, as the constructible/blocked boundary.

The AKS validity twist pins the construction to a sign-split cell: a rigid-cell n can satisfy the component equation of Lemma 3.1(ii) only with exponent $j = 0$, forcing $n \equiv 1 \pmod r$ — inside the exception clause, not a counterexample (cf. Váňa's observation [7] that Carmichael-type numbers pass $T(-1, n, r)$ only in such classes). Concretely: Howe's rigid n has $n \equiv 1 \pmod 5$, yet it does not even pass $T(-1, n, 5)$ (Section 6.1): rigidity controls the $p^2 - 1$ part of $\rho_5(p)$ but not the

extra μ_{2r}-torsion — the factor $2r$ in Lemma 3.2's bound — which is exactly the 2-adic bookkeeping the L–P conditions modulo 80 exist to fix.

Remark 6.5 (imports for the negative-result paper). Three direct imports: (1) *vocabulary* — the paper's target objects are Williams numbers in prescribed congruence cells, an established name with a 1977 pedigree; (2) *search floor* — McIntosh–Mitra's 2^{64} bound is the modern version of the Pinch-table emptiness the README cites; (3) *calibration* — Howe's construction of (d) is the nearest success of the AGP machine to our target cell, and the one-sided/sign-split contrast of Remark 6.4 is the structural explanation of why his succeeded where H8/N1 collapsed; any future attempt on the Williams cell must engage that contrast, not the depth.

6.1 Numerical checks

`code/ag3_hierarchy_check.py` (deterministic; output in `data/generated/ag3_hierarchy_check.json`):

check	statement	result
H1	Howe's n is squarefree, order 2 and rigid; not Williams	confirmed
H2	dictionary on T_2: congruence $\iff$ point identity, random (p, n, i)	60/60
H3	census below 10^7: Carmichael / Lucas–C / Williams / order-2	105 / 135 / 0 / 0
H4	Howe's rigid n passes $T(-1, n, 5)$	**no**

In H1 the matching level-2 powers at every p are $\{0, 2, 4\}$ — all even, the rigid parity; no prime is in the Lucas cell, placing n outside Williams as expected. H3's Carmichael count (105 below 10^7) matches the known value, validating the census code; the smallest members found are 561 (Carmichael) and 399 (Lucas–Carmichael). H4 realizes Remark 6.4's last point: even within the exception clause, rigid impersonation alone does not pass the AKS congruence — the μ_{2r} torsion must be arranged separately.

7 The elliptic axis: formulation and the stabilizer-index floor

Everything measured in this project (A1–A11, N1) lives on $\mathbb{G}_m$ and its tori. Item AG4 asks whether replacing $\mu_r \subset \mathbb{G}_m$ by elliptic torsion opens a route the torus negative results do not cover. The answer, reached here through stages (a) and (b) of the plan, is no: the burden floor of 2 is a structural invariant, and gate AG-G2 is not passed.

7.1 Formulation (AG4a): the test needs CM

The AKS congruence rests on the *canonical* substitution $\sigma_n : X \mapsto X^n$ — the Adams operation of the Λ-ring $\mathbb{Z}[\mu_r]$ (Section 1). If an elliptic curve carried such a canonical commuting family, each ψ_p would be an endomorphism of degree p lifting Frobenius, and an endomorphism of non-square degree already forces $\mathrm{End}(E) \supsetneq \mathbb{Z}$, i.e. complex multiplication — so only CM curves are eligible; conversely, for CM curves the canonical family exists (Borger–Gurney [23]). For a non-CM curve there is no analogue of σ_n at all, and testing against all candidate traces $a \in [-2\sqrt{n}, 2\sqrt{n}]$ is exponential. So the well-posed analogue is CM:

Definition 7.1 (elliptic AKS test, CM form). *Let $E/\mathbb{Q}$ have CM by O_K, K imaginary quadratic with unit group $O_K^\times$ of order w ($w = 4$: $j = 1728$, $E : y^2 = x^3 - x$, $[i](x, y) = (-x, iy)$; $w = 6$: $j = 0$). Fix an odd prime r of good reduction split in K, a kernel polynomial g of an r-eigenline*

(certificate-style, or explicit for special sequences as in [10]), and a section f (we take $f = x$). The test: for composite-candidate n coprime to $2r\,\Delta_E$,

$$\exists\,\tau \in (\mathbb{Z}/r)^\times : \qquad f^{\,n} \;\equiv\; f \circ [\tau] \pmod{n,\; g(x)},$$

$[\tau]$ acting through division polynomials ($r-1$ candidates, so poly(r) trials). For n prime, ordinary and split, the n-power Frobenius acts on the eigenline by a scalar λ_n and the test passes with $\tau = \lambda_n$.

Remark 7.2 (the exception clause). Let $U \le (\mathbb{Z}/r)^\times$ be the image of $O_K^\times$ (for $j = 1728$: $U = \mu_4 = \{\pm 1, \pm i\}$, present iff $4 \mid r-1$). Cells with $\lambda_p \in U$ at every $p \mid n$ are the elliptic analogue of the rigid cells of Section 6: their assembly is a one-sided AGP-type problem (heuristically constructible, as Howe's example demonstrates for the rigid torus cell), so a credible elliptic Agrawal conjecture must read: *the congruence forces n prime or $\tau \in U$* — the clause is the stabilizer, exactly as $n^2 \equiv 1 \pmod r$ is $\sigma_n \in \{\pm 1\}$ (Lemma 1.3). Validity of a counterexample therefore means λ-data escaping U.

7.2 The stabilizer-index theorem (AG4b)

The collapse identities are exact, as in Lemma 3.2: for an eigenline point T with Frobenius scalar λ, $x(T)^{p^e} = x([\lambda^e]T)$, and on $E : y^2 = x^3 - x$ the automorphism gives $x([i]S) = -x(S)$ for *all* S. Hence $\lambda^e \equiv \pm 1$ gives $x(T)^{p^e} = x(T)$ and $\lambda^e \equiv \pm i$ gives $x(T)^{p^e} = -x(T)$; with

$$k' \;:=\; \min\{\, e \ge 1 : \lambda^e \in U \,\} \;=\; [\langle\lambda\rangle : \langle\lambda\rangle \cap U],$$

the order of $x(T)$ divides $2(p^{k'} - 1)$: support on the levels $e \mid k'$, burden k'. The same computation in each earlier setting yields one statement:

Theorem 7.3 (stabilizer-index floor)**.** *Let a test section f on a torsion scheme carry a scalar Frobenius action λ in a cyclic group Γ, and let $U \le \Gamma$ be the subgroup whose action fixes f up to root-of-unity cofactors (the stabilizer of* div f *up to torsion units). Then the p-dependent support of the control order is the levels $e \mid k'$ with $k' = [\langle\lambda\rangle : \langle\lambda\rangle \cap U]$ — an unconditional upper bound via the exact cofactor identities, attained generically in the sense of Proposition 3.5. A valid counterexample requires $\lambda \notin U$ (Remark 7.2 and its torus analogues), hence $\langle\lambda\rangle \cap U \lneq \langle\lambda\rangle$ and*

$$\textit{burden} \;=\; k' \;\ge\; 2.$$

Instances: $(\mu_r,\, X-1)$: $U = \{\pm 1\}$, burden $d/\gcd(d,2)$ adjusted to d' — Theorem 3.6; $(\mu_r,\, X+2)$: $U = \{1\}$, burden d — Theorem 5.5; $(E_{j=1728},\, x)$: $U = \mu_4$, burden $m/\gcd(m,4)$ for $\mathrm{ord}(\lambda) = m$; $(E_{j=0},\, x)$: $U = \mu_6$-image, burden $m/\gcd(m,6)$.

Proof. Collapse: with $e = k'$ minimal, $\lambda^{k'} \in U$ gives $f(T)^{p^{k'}} = \zeta f(T)$ with ζ a root of unity of small order (the cofactor of the unit action on f: $-\zeta^{-1}$ for $X-1$ under inversion, -1 for x under $[i]$, $\zeta_3^{\pm 1}$ for x under $[\zeta_3]$), so $\mathrm{ord}\, f(T)$ divides (small) $\cdot\,(p^{k'} - 1)$. No smaller e works: for $e < k'$, $\lambda^e \notin U$ and the cofactor $f \circ [\lambda^e]/f$ is a non-torsion unit (cyclotomic for μ_r, Lemma 5.3; elliptic for E, where the division values are supported above r and distinct points give distinct prime support — the same disjoint-support argument), so no identity collapses the order; generic sharpness then puts the full levels $e \mid k'$ in the order. The index formula is cyclicity of $\langle\lambda\rangle$. Validity excludes $\lambda \in U$, so the index is at least 2. □

Remark 7.4 (supersingular cells). For p inert in K (e.g. $p \equiv 3 \bmod 4$ on $y^2 = x^3 - x$), $a_p = 0$ and $\varphi^2 = [-p]$, so the collapse of x-values is governed by $\langle p\rangle \subseteq (\mathbb{Z}/r)^\times$ exactly as in the torus case — supersingular cells offer no separate escape.

7.3 Gate AG-G2: not passed

Corollary 7.5 (the elliptic axis is closed)**.** *On the $j = 1728$ curve the per-modulus burden is* $\mathrm{burden}_{\mathrm{ell}}(r) = \min\{\, m/\gcd(m,4) \geq 2 : m \mid r-1 \,\}$*, degenerate (no valid cell) when $r-1 \mid 4 \cdot 2^{\infty}$ fails to provide one, e.g. $r = 5$; on $j = 0$ replace 4 by 6. In every tested case (Section 7.4)* $\mathrm{burden}_{\mathrm{ell}}(r) \geq \mathrm{burden}_{\mathrm{Agr}}(r)$*, and the floor 2 — attained iff $8 \mid r-1$ ($j = 1728$) or $12 \mid r-1$ ($j = 0$) — lands on the levels $\{1,2\}$: $p-1$ and $p+1$ smooth, the same twin-smooth demand as the torus route, now with a ray-class slice (λ_p of order $2w$ mod r, a positive-density Chebotarev condition) in place of the mod-80 cell, and with unit-twisted mixed targets ($x(T)^{p^{k'}} = -x(T)$), i.e. the sign-split structure of Remark 6.4 — so the N1/A10 harvest-and-solver obstruction transfers verbatim. The apparent burden-1 cells ($\lambda \in U$) are exactly the exception clause. Gate AG-G2 of the plan is therefore* ***not passed****: the elliptic axis offers no single-rare-event escape, and stage AG4(c) is cancelled; this section extends the negative result to the CM-elliptic axis.*

Caveats, beyond the standing ones (mechanism scope, Remark 2.3; generic sharpness, Remark 3.9): the formulation itself is CM-bound — the non-CM "elliptic Agrawal" is not well-posed for lack of a canonical Frobenius lift [23] — and the eigenline kernel polynomial is certificate-style data for general n (explicit only for special sequences, as in [10, 24]).

7.4 Numerical validation

`code/ag4_elliptic_check.py` (deterministic; output in `data/generated/ag4_elliptic_check.json`), on $E : y^2 = x^3 - x$ with $r \in \{13, 17\}$ and ten ordinary primes $p \equiv 1 \pmod 4$ each:

check	statement	result
E1	$\deg \psi_r = (r^2-1)/2$; factor degrees multiply back	all
E2	Frobenius closure $t^{p^k} = t$ per factor field	all
E3	exceptional cells ($k' = 1$, $\lambda \in \mu_4$) have $k \leq 2$	45/45
E3	sharpness: large primes of levels $e \mid k'$ divide $\mathrm{ord}(t)$	112/112
E4	measured valid-cell floor: $r = 13 \to 3$, $r = 17 \to 2$	matches index formula
E5	$\mathrm{burden}_{\mathrm{ell}}(r) \geq \mathrm{burden}_{\mathrm{Agr}}(r)$ over the A11 moduli	11/11

The $r = 17$ data shows the floor concretely: degree-4 factors with $t^{p^2} = -t$ (order-8 λ: burden 2, $\Phi_4(p)$ absent from $\mathrm{ord}\, t$) alongside degree-8 factors with $k' = 4$ (order-16 λ: levels $\{1,2,4\}$) and degree-16 mixed cells; at $r = 13$ no valid cell beats burden 3, as $12/\gcd(12,4) = 3$ predicts. $r = 5$ is degenerate on this curve: every λ-order divides 4, so all cells are exceptional.

A The function-field mirror (item AG6, sketch)

Over $A = \mathbb{F}_q[t]$ every object of this document has a classical mirror, and the comparison isolates what our two heuristic caveats — generic sharpness (Proposition 3.5) and the L–P-style density counts — are standing in for over $\mathbb{Z}$.

The mirror of $\mathbb{G}_m$ with its Adams operations is the *Carlitz module* $C : A \to \mathrm{End}(\mathbb{G}_a)$, $C_t(x) = tx + x^q$: the maps C_a commute, and for a monic irreducible P of degree d one has the Frobenius-lift congruence $C_P(x) \equiv x^{q^d} \pmod P$ — here a *theorem* of the module rather than a test. The mirror of $\mathbb{Z}[\mu_r]$ is the Carlitz cyclotomic ring $A[x]/(C_R(x))$ for an auxiliary irreducible R, whose torsion values generate the cyclotomic function fields (Carlitz–Hayes theory). The mechanism analysis of Sections 2–3 transposes: residue fields of torsion values are constant-field extensions $\mathbb{F}_{q^{de}}$, their unit

orders factor through the same cyclotomic levels, and the stabilizer-index principle of Theorem 7.3 applies with U the image of the constants $\mathbb{F}_q^\times$, which act on the identity section with monomial cofactors ($C_c(x) = cx$).

Two ingredients change epistemic status. (i) The sharpness questions — which large primes divide the order of the reduction of a fixed torsion value as P varies — become effective: Frobenius equidistribution in the relevant cyclotomic-function-field extensions holds with square-root error by Weil's RH for curves (function-field Chebotarev), which is exactly the Artin-type input that keeps Proposition 3.5 heuristic over $\mathbb{Z}$. (ii) The existence side (primes P in prescribed cells with smooth order data) becomes quantitative for the same reason. Yet the mirror carries no complexity content: irreducibility testing in $\mathbb{F}_q[t]$ is unconditionally deterministic polynomial time (Frobenius linear algebra: $x^{q^n} \equiv x$ and gcd conditions), so no Agrawal-type conjecture is needed there. The conclusion of the comparison is a placement, not a plan: *the mechanism analysis of this document is complete in itself, and the part of the Agrawal problem that resists it over $\mathbb{Z}$ is exactly the unconditional-equidistribution gap that Weil closes over $\mathbb{F}_q[t]$* — a GRH-grade input, consistent with the fact that the best unconditional test [3] works by constructing its Frobenius datum rather than certifying it from one section. We make no claims here about Carlitz analogues of the Carmichael hierarchy of Section 6; that literature is not needed for the placement above.

References

[1] M. Agrawal, N. Kayal, N. Saxena, PRIMES is in P, *Ann. of Math.* 160 (2004), 781–793.

[2] H. W. Lenstra Jr., C. Pomerance, Remarks on Agrawal's conjecture, AIM workshop notes, 2003.

[3] H. W. Lenstra Jr., C. Pomerance, Primality testing with Gaussian periods, *J. Eur. Math. Soc.* 21 (2019), 1229–1269.

[4] R. Popovych, A note on Agrawal conjecture, Cryptology ePrint 2009/008.

[5] R. Popovych, Elements of high order in finite fields of the form $\mathbb{F}_q[x]/\Phi_r(x)$, *Finite Fields Appl.* 18 (2012), 700–710.

[6] A. Hegde, P. Devaraj, Heuristics for the construction of counterexamples to the Agrawal conjecture, Springer PROMS, 2021.

[7] T. Váňa, Agrawal's conjecture and Carmichael numbers, Comenius Univ. Bratislava, 2009.

[8] B. Kunyavskii et al., Algebraic tori — thirty years after, arXiv:0712.4061.

[9] K. Rubin, A. Silverberg, Torus-based cryptography, CRYPTO 2003.

[10] A. Gurevich, B. Kunyavskii, Deterministic primality tests based on tori and elliptic curves, *Finite Fields Appl.* 18 (2012), 222–236.

[11] J. Borger, Λ-rings and the field with one element, arXiv:0906.3146.

[12] C. Wilkerson, Lambda-rings, binomial domains, and vector bundles over $\mathbb{CP}(\infty)$, *Comm. Algebra* 10 (1982), 311–328.

[13] W. R. Alford, A. Granville, C. Pomerance, There are infinitely many Carmichael numbers, *Ann. of Math.* 139 (1994), 703–722.

[14] H. C. Williams, On numbers analogous to the Carmichael numbers, *Canad. Math. Bull.* 20 (1977), 133–143.

[15] O. Echi, Williams numbers, *C. R. Math. Acad. Sci. Soc. R. Can.* 29 (2007), 41–47.

[16] R. J. McIntosh, D. Mitra, Carmichael numbers with $p + 1 \mid n + 1$, *J. Number Theory* 147 (2015), 81–91.

[17] E. W. Howe, Higher-order Carmichael numbers, *Math. Comp.* 69 (2000), 1711–1719.

[18] O. Eterevsky, M. Vsemirnov, On the number of prime divisors of higher-order Carmichael numbers, *Fibonacci Quart.* 42 (2004), 141–148.

[19] T. Wright, There are infinitely many elliptic Carmichael numbers, *Bull. Lond. Math. Soc.* 50 (2018), 791–800.

[20] J. H. Silverman, Elliptic Carmichael numbers and elliptic Korselt criteria, *Acta Arith.* 155 (2012), 233–246.

[21] A. Ekstrom, C. Pomerance, D. S. Thakur, Infinitude of elliptic Carmichael numbers, *J. Aust. Math. Soc.* 92 (2012), 45–60.

[22] J. Grantham, Frobenius pseudoprimes, *Math. Comp.* 70 (2001), 873–891.

[23] J. Borger, L. Gurney, Canonical lifts and Λ-structures (elliptic curves with complex multiplication), arXiv:1710.08674.

[24] J.-M. Couveignes, R. Lercier, Elliptic periods and primality proving, arXiv:0810.2853.